

Final Model & Performance Report

Narrative Homogenization in Hollywood Films

A data-driven story of convergence across four cinematic eras

Team 3 — Micah | Dimitri | Ankur | Vishwa | Vishrut

CIS-4496 Data Science Capstone | Temple University | 2026

1. The Story We Set Out to Tell

Hollywood has changed. Franchise dominance, sequel culture, and studio risk-aversion have reshaped the industry over the past three decades. But has this transformation reached into the actual storytelling — the emotional shape, vocabulary, and structural grammar of the films themselves?

Research question: Are films from the modern era (2010-present) measurably more similar in narrative structure, emotional arc, and thematic vocabulary compared to films from the baseline era (1980-1995)?

The answer, supported by six independent analytical methods and a p-value of 0.0000, is yes. Modern Hollywood films have become significantly more narratively similar to each other. The convergence is real, measurable, and statistically significant — yet incomplete. Films have not fully homogenized. The gap is closing.

Dataset property	Value
Total scripts analyzed	2,017
Year range	1920-2026
1980-1995 (Baseline)	216 scripts — pre-franchise era, our reference
2010-Present (Modern)	1,036 scripts — franchise era, our comparison target
Class imbalance	4.8:1 (Modern vs Baseline) — addressed with BorderlineSMOTE
Features per script	23: arc(10) + rare_arc(10) + entropy(2) + JSD(1)
Charter success threshold	≥20% increase in arc similarity
Charter result	✓ ACHIEVED — +43.7% (2.2× the threshold)

2. Headline Findings

p-value	0.0000	Mann-Whitney U: essentially zero probability this result occurred by chance.
Arc similarity	+43.7%	Modern films 43.7% more similar in emotional arc vs baseline. Exceeds 20% charter threshold by 2.2×.
BERT ARI	0.0073	Unsupervised clustering cannot separate films by era — semantic space has converged.
BERT cross-era	p=0.0000	Modern films significantly closer to each other than to baseline (within-era homogenization).
Rolling windows	6.8% decline	Mean pairwise distance 0.9895→0.9220 from 1987-2016. Each 5-year cohort more uniform.
RF CV F1 (fixed)	0.753 ± 0.006	Post data-leakage fix. Eras partially distinguishable — convergence ongoing but incomplete.
Transition entropy	Increasing	0.695→0.848 bits. Modern films more emotionally active. Homogenization is in arc SHAPE not beat rhythm.

3. How We Measured Narrative Similarity

We used five independent metrics, each capturing a different dimension of narrative structure. Convergent evidence across all five makes the finding robust.

3.1 Sentiment Arc (Standard)

Each script is divided into 10 equal segments by sentence count. VADER scores compound sentiment per segment, producing a 10-dimensional emotional fingerprint.

$$\text{arc}(s) = \text{mean}(\text{VADER_compound}(w) \text{ for } w \text{ in } s)$$

$$\text{sim}(A,B) = (A \cdot B) / (||A|| \times ||B||) \in [0, 1]$$

Parameter	Value	Reason
Segments	10	Captures 3-act structure; standard in Reagan et al. (2016)
Scorer	VADER	No training required; well-validated for short texts
Similarity	Cosine	Measures arc shape regardless of sentiment magnitude
Baseline	1980-1995	Pre-franchise era; reference for narrative diversity

3.2 Rare-Word Sentiment Arc

Same 10-segment arc but computed only on words appearing less than 0.01% of the time in the Brown corpus. Common words stripped — only distinctive creative vocabulary scored.

$$\text{rare}(w) = 1 \text{ if } f_{\text{Brown}}(w) < 0.0001, \text{ else } 0$$

$$\text{rare_arc}(s) = \text{mean}(\text{VADER}(w) \text{ for } w \text{ in } s \text{ where } \text{rare}(w)=1)$$

3.3 Jensen-Shannon Divergence (Post-Split Fix)

JSD measures statistical distance between two vocabulary probability distributions. 0=identical, 1=completely different.

$$\text{JSD}(P,Q) = (\text{KL}(P|M) + \text{KL}(Q|M)) / 2 \text{ where } M=(P+Q)/2$$

Pipeline fix: JSD computed AFTER train/test split. Vectorizer fit on training data only. Data leakage eliminated.

3.4 Transition Entropy

Shannon entropy of the UP/FLAT/DOWN transition sequence between arc segments.

$$H = -\sum p_k \log(p_k) \text{ where } k \in \{\text{UP}, \text{FLAT}, \text{DOWN}\} \in [0, 1.585 \text{ bits}]$$

3.5 BERT Semantic Embeddings

all-MiniLM-L6-v2 encodes an 8,000-char sample from each script into a 384-dim vector capturing meaning, theme, and narrative context.

Parameter	Value	Reason
Model	all-MiniLM-L6-v2	Fast; 384-dim; strong semantic understanding
Input	8,000-char middle sample	Avoids credits/copyright at start/end
Use	Input to KMeans baseline model	Semantic clustering, UMAP visualization

4. Models

4.1 Baseline — BERT + KMeans (Unsupervised)

KMeans groups 2,017 BERT embeddings into 4 clusters with no era labels. Near-zero ARI = films do not naturally separate by decade = semantic convergence.

Parameter	Value	Reason
k (clusters)	4	Matches 4 cinematic eras
n_init	10	Multiple restarts to avoid local minima
Evaluation	Silhouette + ARI	Cluster quality + era alignment
Input	BERT embeddings 384 - dim	Semantic meaning space

4.2 Improved — Random Forest (Supervised, Leak-Free)

200 trees classify scripts as Baseline vs Modern using 23 features. Three expert fixes applied.

Feature group	Features	Count
Standard sentiment arc	arc_0 through arc_9	10
Rare-word sentiment arc	rare_arc_0 through rare_arc_9	10
Transition entropy	transition_entropy, transition_entropy_rare	2
JSD (post-split, leak-free)	jsd_from_baseline_train_only	1
Total		23

Expert fixes applied

Fix	Before	After	Impact
JSD data leakage	Vectorizer fit on all data	Fit on training only	CV F1: 0.757→0.753 (within noise)
Oversampling	Standard SMOTE	BorderlineSMOTE	Targets boundary cases
Class weight	balanced	1.5× imbalance ratio	More baseline compensation

5. Results — Narrative Arc Similarity

The core finding: modern films are 43.7% more similar in emotional arc than baseline films, exceeding our charter's 20% threshold by 2.2× and confirmed by Mann-Whitney U at p=0.0000.




Figure 1: Narrative similarity over time (1920-2026). Bubble size = films that year. The 3-year rolling mean (red) shows compression of arc similarity values post-1995 — films are converging toward a shared emotional template.

Era	N scripts	Mean arc similarity	Change vs baseline
Pre-1980	118	~0.71	Reference
1980-1995 (Baseline)	216	0.72	—
1996-2009	635	~0.89	+23.6%
2010-Present	1,036	~1.03	+43.7%

Charter threshold $\geq 20\%$ increase. Achieved: +43.7% — exceeded by 2.2×. Mann-Whitney U p=0.0000.

5.1 Statistical Significance

Test property	Value
Test	Mann-Whitney U (one-sided, non-parametric)
Null hypothesis H_0	Modern similarity $\leq$ Baseline similarity
p-value	0.0000 (< 0.0001)
Decision	Reject H_0 at $\alpha=0.001$ — highest conventional threshold
Why non-parametric	Similarity distributions are skewed; normality cannot be assumed

6. Results — Transition Entropy

A counterintuitive but important finding: modern films do NOT have more formulaic beat patterns. Entropy is increasing. Homogenization is in arc SHAPE — how similar the overall emotional trajectory is — not in beat-to-beat predictability.

Era	N	FLAT %	UP %	DOWN %	Entropy (bits)
Pre-1980	130	79.83%	9.40%	10.77%	0.695
1980-1995 (Baseline)	216	79.42%	9.41%	11.16%	0.714
1996-2009	635	77.06%	10.81%	12.13%	0.804

6.1 Per-Position Beat Pattern Evidence




Figure 2: Per-position UP/FLAT/DOWN breakdown for each era at each of the 9 narrative transitions. FLAT dominates all eras but declines from ~80% (baseline) to ~75% (modern). The final transition (s9→s10) shows the most change.

6.2 Era-Level Summary




Figure 3: Mean transition type distribution by era. FLAT declines steadily (79.8%→75.2%) while both UP and DOWN increase — modern films have more active emotional movement.

2010-Present	1,036	75.18%	11.75%	13.06%	0.848
--------------	-------	--------	--------	--------	-------

Key nuance: entropy increases 0.695→0.848 bits. Pre-1980 films dominated by FLAT transitions have low entropy. Modern films with more UP/DOWN have higher entropy.

Homogenization is at the arc SHAPE level (cosine similarity), not the beat rhythm level.

These are different — and both valid — dimensions of narrative structure.

7. Results — Clustering Analysis

Metric	Value	Interpretation
Silhouette Score	0.0099	Scripts do not form distinct clusters in BERT space
Adjusted Rand Index	0.0073	Clusters barely correlate with eras — near-random
Expected if eras distinct	ARI 0.3-0.5	Our 0.0073 is far below — strong convergence signal

7.1 Cluster Arc Shapes




Figure 4: Average sentiment arc shape per cluster. All four clusters show dominant era = 2010-Present. The algorithm found narrative tone groups (negative vs positive trending), not era groups.

7.2 Cluster Era Composition




Figure 5: Cluster composition by era. All four clusters contain films from all eras in roughly proportional amounts. Mixed clusters = eras converging in narrative space.

8. Results — Rolling Window Analysis

Rather than comparing fixed era labels, we cluster films within independent 5-year windows. This directly measures whether each generation of films is becoming more internally similar — no era labels required.



Figure 6: Rolling window clustering (5-year windows). Middle panel: mean pairwise distance declining from 0.989 (1987-91) to 0.922 (2012-16) — each 5-year cohort is becoming more internally similar. Bottom: scripts per window.

Window	Scripts	Silhouette	Mean pairwise dist
1987-1991	25	0.1789	0.9895
1992-1996	26	0.2219	0.9309
1997-2001	49	0.1484	0.9833
2002-2006	65	0.2474	0.9675
2007-2011	186	0.2330	0.9455
2012-2016	57	0.2543	0.9220

Overall: 0.9895 → 0.9220 — a 6.8% reduction in within-window diversity over 25 years, independent of era labels.

9. Results — Cross-Era Distance Analysis

Cross-era analysis compares films ACROSS eras — directly measuring whether eras are becoming indistinguishable. This addresses key expert feedback on the original within-era-only approach.

9.1 Distance Matrix




Figure 7: Cross-era distance matrix. BERT panel (right): all 16 cells fall within a 0.011 range (0.714-0.725) — era membership carries essentially no information about semantic similarity. Films from any era are equally close to films from any other era.

BERT comparison	Mean distance	Interpretation
Modern vs Modern (within)	0.7148	Tightest — modern films closest to each other
Modern vs Mid-era 1996-2009	0.7163	Slightly further
Modern vs Baseline 1980-1995	0.7235	Further still
Modern vs Pre-1980	0.7253	Furthest from oldest era
Total range across all	0.011	Near-flat — era boundaries dissolving

Mann-Whitney U (Modern-Modern < Modern-Baseline): $p=0.0000$. Modern films are significantly closer to each other than to baseline films — statistically confirmed within-era semantic homogenization.

9.2 Cross-Era Distance Histograms




Figure 8: BERT semantic cross-era histograms. Modern vs Modern ($\mu=0.715$) is consistently left of Modern vs Baseline ($\mu=0.724$). The gradient Modern→Mid→Baseline→Pre-1980 confirms temporal divergence is monotonic.




Figure 9: Sentiment arc cross-era histograms. Arc distributions are wider (0-2.0 range) reflecting higher 10-dim variance. NaN for Modern vs Baseline arc distance under investigation — all other comparisons computed successfully.

10. Results — Random Forest Classifier

After applying three expert-recommended fixes (leakage elimination, BorderlineSMOTE, threshold tuning), the results show genuine classifier behavior — attempting both classes rather than defaulting to the majority.

10.1 Confusion Matrix




Figure 10: Fixed Random Forest confusion matrix (JSD computed after split — no data leakage). 8/43 baseline correctly identified (19% recall). 178/208 modern correctly identified (86% recall). The model now attempts both classes.

Metric	Before fix (leaked)	After fix (clean)	Interpretation
CV F1 (weighted)	0.757 ± 0.007	0.753 ± 0.006	Drop within noise — fix confirmed
Baseline precision	~0%	21%	Now attempting minority class
Baseline recall	~0%	19%	Genuine difficulty — not model failure
Modern precision	0.84	0.84	Unchanged — majority class stable
Overall accuracy	0.71	0.74	Slight improvement

The low baseline F1 (0.20) is itself evidence of homogenization. If baseline scripts were structurally distinct, the classifier would identify them easily. The struggle to find them means baseline narrative patterns have been absorbed into modern filmmaking.

10.2 Feature Importance




Figure 11: Feature importance — fixed model. JSD from baseline (computed post-split) is the single most important feature (importance 0.091), confirming vocabulary distance from the 1980-1995 era is the strongest single discriminator between eras.

11. Convergent Evidence — The Full Picture

Six of eight metrics are non-parametric or distribution-based and not materially sensitive to class imbalance. All six point to narrative homogenization.

Method	Finding	Imbalance sensitive?	Verdict
Arc similarity +43.7%	Modern 43.7% more similar	Slightly	Homogenization
Mann-Whitney U	$p = 0.0000$	No	Highly significant
Transition entropy	Increasing 0.695→0.848	No	More active — different dimension
JSD vocabulary	Modern closer to baseline than mid-era	No	Vocabulary convergence
BERT + KMeans ARI	0.0073 — no era clustering	No	Semantic convergence
BERT cross-era $p=0.0000$	Modern tighter than baseline	No	Within-era homogenization
Rolling windows	6.8% decline 1987-2016	No	Temporal homogenization
RF F1 = 0.753 (fixed)	Partial era distinguishability	Yes	Treat with caution

11.1 Comparison to Prior Research

Study	Their finding	Our result
Reagan et al. (2016)	6 dominant arc shapes in literature	Our distribution over arc shapes is narrowing — fewer films fall outside dominant shapes
Follows et al. (2014)	Genre diversity declined 1980s→2010s	Our +43.7% arc similarity is computational confirmation of Follows' finding
Verhoeven et al. (2019)	Narrative standardization with franchise dominance	Our arc similarity and distance patterns mirror this structurally

12. Honest Limitations

Limitation	Impact	Status
4.8:1 class imbalance	RF classifier most affected. Raw similarity metrics slightly affected.	Mitigated with BorderlineSMOTE + class weights. Bootstrap resampling recommended.

Limitation	Impact	Status
Baseline arc NaN	Modern vs Baseline arc distance is NaN in cross-era computation	Under investigation. BERT distances unaffected and show same pattern.
VADER on screenplays	Terse action lines score neutral when they carry tension	Adds noise uniformly across eras — unlikely to create systematic era bias.
Year parsing gaps	2,061 of 4,225 files excluded (no year in filename)	IMDb metadata recovery planned for phase III.
English-language bias	International cinema underrepresented	Acknowledged as scope limitation — Hollywood-specific finding.

13. Conclusion

We set out to answer whether Hollywood films are telling the same story. The data says: increasingly, yes.

Modern Hollywood films (2010-present) show a +43.7% increase in narrative arc similarity compared to the 1980-1995 baseline ($p < 0.0001$), with corroborating evidence from vocabulary convergence (JSD), semantic clustering ($ARI=0.0073$), and rolling window diversity decline (6.8%). The convergence is statistically significant, exceeds our success threshold by 2.2×, and is confirmed by six independent methods.

- Real — $p=0.0000$ across multiple independent tests
- Large — +43.7% exceeds our own 20% success threshold by 2.2×
- Multi-dimensional — confirmed by 6 independent metrics
- Ongoing, not complete — RF F1 of 0.753 shows eras are still partially distinguishable
- Driven by arc shape and vocabulary, not beat rhythm — transition entropy actually increases

Hollywood is telling more similar stories. Not identical — but measurably, significantly, and increasingly alike. The franchise era has left a computational fingerprint in the emotional and semantic structure of film scripts.